
Stage Four: Sprint Plan

Author: Madhav Mehta

Section one: capacity

An individual project alongside other commitments. Realistic rather than optimistic capacity, because a plan built on best-case hours produces a backlog that is abandoned in week two.

Week	Available hours	Committed to	Contingency held
One	28	Discovery, requirements, environment	4 h
Two	34	Build, evaluation, revision	6 h
Three	26	Governance, report, video, packaging	5 h
Total	88		15 h (17%)

Contingency is held rather than allocated. Every previous estimate I have made on a project of this shape has been optimistic by roughly a fifth, and the effort log records whether that held again.

Section two: the backlog

Owner is the author throughout; the column is retained because the workbook asks for it and because "definition of done" is the part that actually matters.

ID	Item	Est	Actual	Depends on	Definition of done
T-01	Environment, dependency pins, repository skeleton	3 h	2 h	—	<code>pip install -r requirements.txt</code> succeeds from an empty venv
T-02	Read all five transcripts; note claims to verify	3 h	4 h	—	Every claim that can be checked against data is written down as a question
T-03	Discovery analysis script	4 h	6 h	T-01, T-02	One command regenerates every figure in the workbook
T-04	Stage 1 workbook	5 h	6 h	T-03	Every row cites either a transcript or a computed figure
T-05	PRD v1	4 h	4 h	T-04	Every requirement names its discovery evidence
T-06	Ingestion, four channels	3 h	2 h	T-01	A ticket from each channel normalises; malformed input does not raise
T-07	Retrieval and chunking	5 h	8 h	T-06	Ids resolve to real articles; an irrelevant query returns nothing
T-08	Intent classifier	4 h	3 h	T-06	

ID	Item	Est	Actual	Depends on	Definition of done
					Per-class precision and recall reported, not just accuracy
T-09	Router and threshold	4 h	9 h	T-07, T-08	Deterministic; threshold justified from data, not chosen
T-10	Generation with citations	4 h	4 h	T-07	Every citation resolves to a retrieved passage
T-11	Guardrails, input and output	4 h	5 h	T-10	A guardrail blocks a real attempt, and does not block an acceptable answer
T-12	Decision log and reconciliation	3 h	4 h	T-09	Logged decisions equal tickets processed, and the check can fail
T-13	Evaluation harness	4 h	4 h	T-06..T-12	--input / --output ; a full unseen file processed with no intervention
T-14	Metrics and report generation	4 h	6 h	T-13	Every pass/fail computed from the value; none written by hand
T-15	Calibration and threshold selection	3 h	7 h	T-08, T-09	Calibration error measured; threshold selected against a stated cost model
T-16	Test suite mapped to A1–A12	5 h	6 h	all	One command; a failure names the criterion it broke
T-17	PRD v2 and revision log	3 h	3 h	T-15	At least one substantive change with its measured trigger
T-18	Governance framework	5 h	5 h	T-12, T-14	Risk register with owners; fairness audit with method and sample size
T-19	README and architecture notes	4 h	5 h	all	Followed literally from an empty directory by someone else
T-20	Report	8 h	9 h	all	Every claim carries how it was measured and on what
T-21	Video	6 h	—	T-19, T-20	18–22 min; live demo including an escalation and a guardrail firing
T-22	Packaging and final checks	3 h	—	all	Four folders, correct names, clean-checkout rehearsal passed

Section three: week two, day by day

Built to the Build Specification checkpoints, with the end-to-end run pulled to day four rather than day five.

Day	Target	Checked by
1	Ingest. A ticket from each channel normalises	Print the normalised object for one ticket per channel
2	Retrieval. Corpus indexed, passages traceable	Query a known question, confirm the right article comes back and a nonsense query returns nothing
3	Classification and routing. Confidence scored, threshold applied, every decision logged with a reason	Run 20 tickets and read the log end to end
4	Full chain over the full validation set, unattended	Start it, leave, return to a metrics report

Day	Target	Checked by
5	Guardrails blocking; calibration; threshold re-derived	An engineered ticket blocks; calibration error measured

Why the run moved to day four. The Build Specification puts it on day five. Moving it one day earlier buys a whole day to react to what it says, and the brief is explicit that the most common failure is discovering results too late to act on them. It was the single best scheduling decision in the project: the day-four run is what surfaced the calibration problem, and had it happened on day five the threshold would have shipped uncorrected.

Section four: week three, day by day

Day	Target
1	Governance: risk register, fairness audit against the run, incident procedure, kill switch
2	CI; clean-checkout rehearsal from an empty directory on a different machine
3	Report sections 1–5
4	Report sections 6–10; first video take
5	Second video take; packaging; final checklist

Section five: what gets dropped if time runs out

Decided in advance, in priority order, so that the decision is not made at midnight on the last day.

Priority	Item	Why it goes in this order
Dropped first	Tier 1 semantic retrieval benchmarking	The fallback is measured and works; the gap is an honest caveat rather than a defect
Second	Prometheus and Grafana monitoring	Nothing in A1–A12 requires it. The decision log already answers the operational questions
Third	Depth in the fairness audit beyond the three segments	Three segments satisfy the requirement; more would be better but is not load-bearing
Fourth	A second video take	Costs polish, not correctness
Never dropped	The unattended run, the decision log, the guardrails, the honest metrics	These are the gate and the difference between a working system and a demonstration

In the event, the first three were dropped. Monitoring and the semantic benchmark are both recorded as outstanding work rather than quietly omitted.

Section six: daily check-in record

Day	Planned	Actual	Carried
W1 D1	Environment, read the brief	Done, plus transcripts	—

Day	Planned	Actual	Carried
W1 D2	Transcripts and analysis script	Script took longer; the by-intent answerability breakdown was worth the extra time	2 h
W1 D3	Stage 1 workbook	Done. The fluency finding contradicted Sofia and needed rechecking	1 h
W1 D4–5	PRD v1	Done	—
W2 D1	Ingest, retrieval	Ingest quick, retrieval slow — chunking strategy reworked twice	3 h
W2 D2	Classifier, router	Classifier quick. 100% accuracy looked wrong and needed investigating	—
W2 D3	Generation, guardrails, logging	Done	—
W2 D4	Full unattended run	Ran. Surfaced the calibration problem.	—
W2 D5	Calibration, threshold, PRD v2	Overran badly; the routing signal was replaced entirely	5 h
W3 D1	Governance	Done. The fairness audit found a 20-point tier gap	—
W3 D2	Tests, clean checkout	The test suite found two real bugs	2 h
W3 D3–5	Report, video, packaging	In progress	—

The one that mattered. W2 D4's run is the reason this project has a genuine revision to report rather than an invented one. Everything that changed in PRD v2 came out of that afternoon.